

Efficient Static Checking of Library Updates

Darius Foo
CA Technologies
Singapore
darius.foo@ca.com

Hendy Chua
CA Technologies
Singapore
hendy.chua@ca.com

Jason Yeo
CA Technologies
Singapore
jason.yeo@ca.com

Ang Ming Yi
CA Technologies
Singapore
mingyi.ang@ca.com

Asankhaya Sharma
CA Technologies
Singapore
asankhaya.sharma@ca.com

ABSTRACT

Software engineering practices have evolved to the point where today, a developer writing a new application doesn't start from scratch, but reuses a number of open source libraries and components. These third-party libraries evolve independently of the applications in which they are used. As bugs and vulnerabilities in these libraries are fixed, they may not be able to maintain stable interfaces. This in turn causes API incompatibilities with downstream applications which must be manually resolved. Oversight here may manifest in many ways, from test failures to crashes at runtime. To address this problem, we present a static analysis for automatically and efficiently checking if a library upgrade introduces an API incompatibility.

Our analysis does not rely on the reported version information from the library developer but instead computes the actual differences between the methods in the library across different versions. The analysis is also scalable, enabling real-time diff queries involving arbitrary pairs of library versions. This in turn enabled us to release a vulnerability remediation product which suggests library upgrades automatically and is lightweight enough to be part of a continuous integration/delivery (CI/CD) pipeline. To demonstrate the effectiveness of our approach, we evaluate semantic versioning adherence on a corpus of open source libraries taken from Maven Central, PyPI, and RubyGems. We find that on average, 26% of library versions are in violation of their semantic versioning scheme. We also analyze a collection of popular open source projects from GitHub to determine if we can automatically update libraries in them without causing API incompatibilities. Our results indicate that we can suggest upgrades automatically for 10% of the libraries.

KEYWORDS

automated remediation, library upgrades, call graphs, api diffs, semantic versioning

ACM Reference Format:

Darius Foo, Hendy Chua, Jason Yeo, Ang Ming Yi, and Asankhaya Sharma. 2018. Efficient Static Checking of Library Updates. In *Proceedings of The 26th ACM Joint European Software Engineering Conference and Symposium on the Foundations of Software Engineering (ESEC/FSE 2018)*. ACM, New York, NY, USA, 6 pages.

ESEC/FSE 2018, 4–9 November, 2018, Lake Buena Vista, Florida, United States 2018.

1 INTRODUCTION

The use of open-source and third-party components has increased in the development of software. Central package distribution systems like Maven Central for Java, RubyGems for Ruby and PyPI for Python make the task of downloading and using these components very easy for the average developer. Unfortunately, these third-party libraries evolve independently of the applications that use them. Bugs and vulnerabilities in third-party components are hard to trace and fix in downstream applications. Even when vulnerabilities can be fixed by updating to a newer version of a library, there can be API incompatibilities with downstream applications which must be manually resolved. Oversight here may manifest in many ways, from test failures to crashes at runtime.

One solution in this space is semantic versioning¹ (SemVer): the idea of adopting a structured versioning scheme that others may rely on. In this form of versioning, the library developers release version numbers in triplets of the form MAJOR.MINOR.PATCH. When introducing a breaking or incompatible API change, developers can release a major version to signal to downstream applications that the changes made are backward-incompatible. As this versioning scheme is structured, it can be consumed by automated tools – an example is Bundler's `~>` operator, which only upgrades packages across patch version boundaries. However, the compliance of *source code* to the scheme must be manually enforced, and it has been criticized as inadequately capturing the nuances of differences in software versions [1].

To offer a better solution, we present a static analysis for automatically and efficiently checking if a library upgrade introduces an API compatibility (colloquially termed a *breaking change*). We do not rely on a surrogate source of truth such as semantic versioning, and instead statically compute the differences between source-level elements of the library (in particular, methods and functions), taking control flow into account. Our analysis is also scalable, enabling real-time diff queries involving arbitrary pairs of library versions.

Our solution forms the basis for a new feature in the CA Veracode Software Composition Analysis product which suggests library upgrades automatically and is lightweight enough to be part of a continuous integration/delivery (CI/CD) pipeline [14]. Left to their own devices, developers do not update dependencies [9], as it is seen as “extra effort and added responsibility”, and the downsides of failing to do so are not as visible. We believe that the ability to have upgrades automatically carried out – but not carelessly, *with*

¹<https://semver.org/>

guarantees about their effects – would go a long way towards eliminating easily-preventable classes of mistakes and vulnerabilities. A recent study on automated pull requests [11] reached the same conclusions, finding that automation caused a 60% increase in upgrades, and that notification fatigue and concerns about breaking changes became the bottleneck thereafter. Our analysis directly address the problem of breaking changes whereby we can statically show the user which library upgrades are possible without introducing API incompatibilities in their application.

To demonstrate its effectiveness, we evaluate semantic versioning adherence on a corpus of open source libraries collected from our customer scan data. These libraries cover three different languages from their respective central repositories – Java (Maven Central), Python (PyPI), and Ruby (RubyGems). We find that on average, 26% of library versions are in violation of their semantic versioning, i.e. they break backward compatibility without a major version update. We also analyze a sample of popular open source GitHub projects to determine the prevalence of API incompatibilities in practice. We find that using our static analysis we can automatically suggest upgrades for 10% of the libraries in these open source projects.

Our main technical contributions are:

- A static analysis that detects breaking changes in libraries accurately, allowing upgrades to be suggested.
- A novel refinement to the analysis which enables diff queries on arbitrary library version pairs to be answered in real time, at the cost of linear space (instead of quadratic).
- A case study of open source libraries published on Maven Central, PyPI, and RubyGems to assess adherence to semantic versioning. On average, 72% of libraries violate SemVer in some version, and 26% of all library versions violate SemVer.

This paper is structured as follows:

- Section 2 compares our techniques with prior attempts to diffs APIs, perform automated library upgrades, and study adherence to semantic versioning.
- In Section 3, we explain how we compute diffs between library APIs, consider control flow, and compose diffs to allow queries on arbitrary version pairs to be answered quickly.
- In Section 4, we detail our study of open source library repositories and evaluate our method of computing diffs on actual projects.
- In Section 5, we discuss areas in which our static analysis may fall short.
- Finally we conclude in section 6 and discuss areas for future improvement.

2 RELATED WORK

2.1 Automated library upgrades

Prior work suggesting automated upgrades ranges from following very simple rules, such as always updating all dependencies within their constraints and relying on test suites to check for breakage², to sophisticated attempts that actually patch APIs or dependent code [8, 18] by inferring new API usage from examples. In contrast,

our approach statically computes diffs to check for breakage and automatically update libraries that do not cause incompatibilities.

Other static analysis approaches for analyzing library upgrades in prior work have made use of dependency analysis [16], symbolic execution [3] and JML contracts [19] to model the semantics of changes.

2.2 API diffs

Textual, subsequence-based diffs, such as those computed by the Unix *diff* and *git* tools, are widely used for comparing program fragments and sharing patches. However, they do not take into account the structure of programs. *Syntactic diffs* may be viewed as a form of diff which considers syntax and is able to ignore textual details. *Syntactic API diffs* include only program elements intended for external use, such as classes and methods. Examples are the documents published by Apple⁴ and Google⁵ to summarize differences between versions of their mobile APIs. UMLDiff [17] and GumTree [6] are implementations of structural diffs that are not specialized to only APIs.

The design space for diffs which reflect *semantic information* is equally wider. SemDiff [4] uses information derived from different subsets of program elements, such as class hierarchies and methods. While Mezzetti et al. [10] uses a dynamic analysis based on the test suites of dependent libraries to infer library interfaces, allowing them to be compared to determine changes. Our approach sits in this category, as we compute syntactic API diffs that are enriched with control flow information, on top of these we apply diff composition to compare arbitrary versions of the library.

2.3 SemVer compliance

We conduct a case study of three open source ecosystems: Maven Central, RubyGems, and PyPI to evaluate the adherence of library developers to SemVer scheme. Prior work in this area by Raemaekers et al. [13] goes into much greater detail, but only for Maven Central. Other related studies evaluate breaking changes in open source projects on npm [10] and CRAN [2].

3 APPROACH

3.1 Basic diffs

Consider the problem of computing diffs for library APIs. We begin by computing a minimal, language-agnostic representation of a library's API, which we term a *signature*. The representation we use is a set of tuples of an *identifier* and a *hash*.

The identifiers give canonical names to the program elements in libraries that we wish to compare. For example, the methods of an object-oriented API, are represented as a tuple of module, class, method names and an argument descriptor.

The hashes summarize the content of the program element, allowing us to quickly determine if it has changed – for example, we hash the bytecode of Java methods, eliding syntactic features such as variable names, but including literals and instructions that affect control flow.

⁴<https://developer.apple.com/library/archive/releasenotes/General/iOS10APIDiffs/index.html>

⁵<https://developer.android.com/sdk/apiff/p-dp1/changes>

```

class A {
    public int a() {
        return 2;
    }
    public int b(int x) {
        return x + 3;
    }
}

```

Figure 1: Example: v1

```

class A {
    // Method a deleted
    public int b(int x) {
        // Modified
        return x + 2;
    }
    // Method c inserted
    public int c() {
        return 1;
    }
}

```

Figure 2: Example: v2

method	operation
A.a()I	DELETE
A.b(I)I	CHANGE
A.c()I	INSERT

Table 1: Example: computed diff

Given the signatures of two libraries, we use Myers’ algorithm [12] to compute a *diff*: an edit script relating them, with the slight modification that we key elements by identifier instead of position in a sequence. This gives us a set of tuples of identifier and *diff operation*, where the latter is one of the symbols **INSERT**, **DELETE**, or **CHANGE**. **CHANGE**s are opaque as program content is summarized using a hash. In particular, we do not perform a more granular tree-based diff like GumTree [6].

Since diffs are meant to relate the *public APIs* of libraries, as a post-processing step we exclude program elements which we know not to be publicly accessible. For a language such as Java, this information is explicit; for more dynamic languages such as Python and Ruby, we rely on heuristics (based on variable names and statically-known calls to methods that modify access) to exclude such elements.

To illustrate all of this, consider two versions of a Java class (Figures 1 and 2). The computed set of diff operations between these two versions is shown in Table 1.

3.2 Transitively-changed methods

Hashing the contents of methods gives us an approximate way to detect changes, but as the above approach hashes methods in isolation and ignores inter-procedural control flow, it does not detect transitive changes to methods. For example, if public method m_1 calls private method m_2 and only m_2 has changed, we would

```

// Before
class A {
    public int m1(int x) {
        return m2(x);
    }
    private int m2(int y) {
        return y + 1;
    }
}

// After
class A {
    // Syntactically unchanged, but
    // returns a different result and
    // so should have a different hash
    public int m1(int x) {
        return m2(x);
    }
    private int m2(int y) {
        return y + 2; // Changed
    }
}

```

Figure 3: Example: transitive changes

miss the fact that m_1 ’s semantics is now different if we exclude m_2 due to its private access. This is illustrated in Figure 3.

To solve this problem, we first build call graphs of library programs. Our call graph construction algorithm uses standard techniques of CHA [5] and VTA [15] and is already available as part of the CA Veracode Software Composition Analysis product. We use the information from the call graphs to improve diffs. Given a method m whose identifier is present in a diff, we also include all public methods that call m . m may or may not be later dropped from the diff, depending on if it is public, but we will no longer exclude its public callers.

3.3 Fast queries

Call graph construction is thus necessary for the accuracy of diffs, but imposes nontrivial overhead. For instance, the largest libraries on Maven Central may take hours to completely analyze. This makes it infeasible to compute accurate diffs on demand, say as part of an automated library upgrade step in a CI/CD pipeline. In this section, we describe a means of precomputing information that enables real-time diff queries on arbitrary pairs of library versions.

A naive approach would be to precompute and store every pair of diffs, consuming space that grows quadratically with the number of versions. This is unlikely to scale since real-world libraries have hundreds of versions, and in general libraries may have up to one version *per commit*⁶. Doing this for a commonly-accessed subset of libraries and a window of recent versions is feasible, but shifts the problem to determining this subset, and of course only works for versions with the subset.

⁶<https://mvnrepository.com/artifact/com.lihaoyi/ammonite-terminal>

The approach we take is to precompute diffs between only *consecutive pairs* of library versions; we then compose individual diffs to derive diffs for arbitrary version ranges. This strikes a good balance, requiring a linear amount of space and running quickly enough in practice.

3.4 Diff composition

We formalize diff composition in this section. As defined earlier, a diff is a set of tuples of a (method) identifier and some *diff operation* which describes how the method changed across versions.

There are 5 primitive diff operations – we saw **INSERT**, **DELETE**, and **CHANGE** earlier, and now introduce two more for explicitness: **UNCHANGED**, which means that a method appears unchanged across versions, and **MISSING**, for when it is missing altogether.

A binary composition operator on diffs must obey two properties: it cannot rely on information external to its operands, and it must be conservative. The latter is the basis for retaining **CHANGE** as a primitive diff operation, instead of expressing it in terms of **INSERT** and **DELETE**: we want to express that a method that is deleted and later inserted might have changed in the process, and not lose information.

Intuitively, **UNCHANGED** and **MISSING** appear to be no-ops. We distinguish them because certain compositions of operations are absurd, e.g. **INSERT** followed by **INSERT**. The underlying reason is that the validity of an operation depends on the state of the method in the versions it acts on: a method must be *absent* for an **INSERT**ion of it to make sense, and it must be *present* thereafter. **UNCHANGED** and **MISSING** operate on methods with different states.

Taking this idea further, we may represent the diff operations collectively as a type indexed by the states of the method they apply to *before* and *after* they are applied. Each diff operation becomes an inhabitant of this type, indexed appropriately.

```
data State = Absent | Present

data Diff : State -> State -> Type where
  Insert : Diff Absent Present
  Change : Diff Present Present
  Delete : Diff Present Absent
  Unchanged : Diff Present Present
  Missing : Diff Absent Absent
```

The composition operation between diffs must then have the following type:

```
compose : Diff a b -> Diff b c -> Diff a c
```

In this way, we partially reduce the problem of checking the validity of a particular composition to determining how the operands of compose constrain its result.

Interestingly, compose is uniquely defined on many inputs. For example, the result of compose Delete Missing *must* be Delete as no other operation is well-typed. compose is also total, as illegal compositions such as compose Insert Insert are not well-typed.

The only ambiguity arises when selecting between **CHANGE** and **UNCHANGED**; as we do not model hashes in our types, they have the same type Diff Present Present (in other words, only

	I	C	D	U	M
I	⊥	I	M	I	⊥
C	⊥	C	D	C	⊥
D	C	⊥	⊥	⊥	D
U	⊥	C	D	U	⊥
M	I	⊥	⊥	⊥	M

Table 2: Diff composition function

	I	C	D	UM
I	⊥	I	UM	I
C	⊥	C	D	C
D	C	⊥	⊥	D
UM	I	C	D	UM

Table 3: Conflated diff composition function

methods which are present throughout may be said to have changed or remained unchanged). We resolve the ambiguity manually, choosing **CHANGE** where possible as it is strictly more conservative than **UNCHANGED**.

The final composition function is given in Table 2. Rows are the first argument and columns are the second. We represent ill-typed combinations with the ⊥ symbol.

Composition is not symmetric:

```
compose Insert Delete = Missing
compose Delete Insert = Change
```

However, it is associative (which can proven by exhaustion).

3.5 Conflating operations

It turns out that we can conflate **UNCHANGED** and **MISSING** into a single operation, **UNKNOWN** (abbreviated **UM**), since they occur in mutually exclusive scenarios. This is useful in practice. Say we diff a list of 100 items against a list of 101; we would want to store a single **INSERT** instead of also storing 100 **UNCHANGED**s. Defaulting to **UNKNOWN** when an item is absent allows us to store the information concisely. This doesn't change composition semantics (proven by exhaustion).

Implementing this change gives us the new function, in Table 3. The time complexity of diff composition is linear in the number of library methods and versions, like diff computation itself.

3.6 Suggesting upgrades

Finally, we use library diffs to suggest upgrades and determine if they induce breaking changes. In the CA Veracode Software Composition Analysis product we already identify vulnerable versions of a library, we now can suggest upgrades to fix those versions. Thus, given a library at version v_1 , and the set of versions v_s of the same library, we choose another version from v_s which succeeds v_1 and does not possess vulnerabilities⁷ associated with v_1 . This may be done using various heuristics and may be subject to further optimization; we currently choose the closest version to minimize diff size and induce the fewest breaking changes.

Given the library diff and the pair of library versions involved in the upgrade l_{from} and l_{to} , we restrict it to only **DELETE** and **CHANGE** operations; this is exactly the set of *affected methods* in

⁷Vulnerability data is assumed to come from an external source, such as NVD or a proprietary vulnerability database [20]


```

16 pom.xml
@@ -17,7 +17,7 @@
17 <dependency>
18 <groupId>org.apache.struts</groupId>
19 <artifactId>struts2-core</artifactId>
20 - <version>2.5.12</version>
20 + <version>2.5.13</version>
21 </dependency>

```

Figure 4: Generated patch to pom.xml

Type	Library	From	To	Breaking
MAVEN	org.apache.struts struts2-core	2.5.12	2.5.13	No
MAVEN	net.bull.javamelody javamelody-core	1.59.0	1.62.0	Possibly

Table 4: Report in generated pull request

l_{from} , in the sense that callers of these methods will be affected by the upgrade.

This information is used when building call graphs of user code to check if they are calling any *affected methods*; if they are, we consider the upgrade to be a breaking change. We use SGL [7], a domain-specific language for program analysis, to implement this portion using a call graph traversal. Details of SGL are described in [7] and are out of scope for this work.

The next step is to generate a patch, rendered in Figure 4 as a GitHub pull request. Since we are patching files which are typically maintained by hand, we take care to minimize changes by parsing the files, then using position information to make granular edits. An example of what we include in the pull request we create is shown in Table 4: we specify the to- and from- versions of the upgrade and say whether or not we were able to statically determine that the change was breaking.

4 EXPERIMENTS AND EVALUATION

4.1 SemVer compliance

We analyzed 5106 libraries (based on customer scan data) across three popular open source library repositories: Maven Central (3273 libraries), RubyGems (1332), and PyPI (501), computing diffs for 114,199 library versions in total.

Our results indicate that on average, 72% of libraries violate SemVer in *some* version: the actual numbers are 80% (RubyGems), 67% (Maven Central), and 82% (PyPI). In addition, on average, 26% of all library versions are in violation, actual numbers being 31% (RubyGems), 24% (Maven Central), and 31% (PyPI). The figure for Maven Central agrees with prior work [13], which puts the number of violating versions between 28.4% and 23.7% over time. The overall distribution of violations is shown in Figure 5, as a plot of the number of libraries (y-axis) that have a given percentage of versions within them in violation (x-axis).

As a concrete example, consider the popular requests library, between versions 2.3.0 and 2.4.0 – a minor version bump, over which breaking changes are not expected – requests.structures.

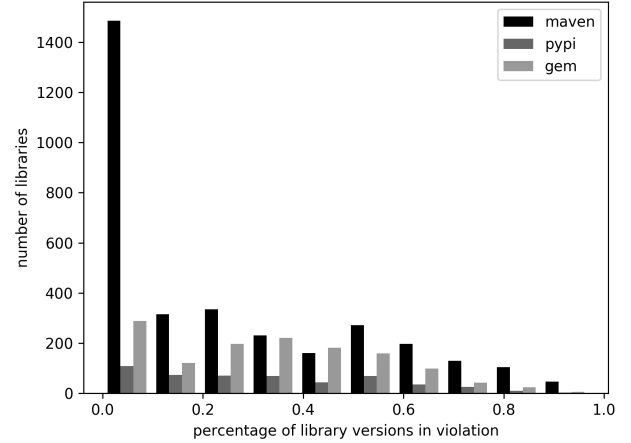


Figure 5: Semantic versioning compliance

	Java	Python	Ruby
Projects	274	422	503
Direct dependencies	4777	2572	4096
Direct vulnerable dependencies	246	110	250
Suggested upgrades	150	64	123
Non-breaking	28 (19%)	0 (0%)	7 (6%)

Table 5: API incompatibilities on GitHub

IteratorProxy was deleted⁸. For this not to be considered a SemVer violation, IteratorProxy must not be part of requests’ public API, however it is difficult to determine this using a static analysis (as nothing prevents one from importing requests.structures). Consulting human-readable sources like commit messages and the change log also yielded nothing in this case, so we assume it to be a breaking change.

4.2 API incompatibilities in open source projects

We analyzed a collection of popular open source projects from GitHub to determine the impact of API incompatibilities on suggesting upgrades automatically. A dynamic dependency analysis was first performed to identify libraries included: e.g. as shown in Table 5, in Java we identified 4777 unique direct dependencies, of which 246 had known vulnerabilities in our database. We were then able to suggest upgrades for 150 of those; 36 had no safe versions to upgrade to, and 60 failed due to the errors in an earlier part of the pipeline (such as malformed class files or the inability to successfully compute a call graph). Of the 150 upgrades, we are able to statically show that 28 (18.7%) are non-breaking.

On average across languages, 10% of upgrades are shown to be non-breaking. The numbers for Python and Ruby are notably lower, and we believe this is due to the difficulties of static call graph construction for those languages, and the false positives that result from over-approximation.

⁸<https://github.com/requests/requests/compare/v2.3.0...v2.4.0#diff-2bdb7e19f5215e8c319573cdd114f01L16>

5 DISCUSSION

5.1 Threats to validity

Limitations of static analysis. A call graph computed using CHA/VTA is an over-approximation of the dynamic control flow of a program. As such, false positives occur, leading to call graph edges that may not actually be traversed in a running program. This may lead to library upgrades being considered breaking when they are not.

Unsupported language features. False negatives also occur, due to lack of support for dynamic language features such as reflection. These manifest as missing edges in the call graph, causing breaking changes to be missed.

Another factor besides dynamism is that program elements interact in more ways than just function and method calls. Java programs are allowed to inspect the class path and behave differently depending on whether certain classes are available; examples of this may be found in real libraries such as the Spring framework. These capabilities allow the semantics of programs to depend on the absence or presence of classes at runtime, something a static analysis would have to approximate and handle in order to accurately determine if an upgrade is breaking.

Computing diffs in isolation. Call graphs and diffs are computed for single libraries at a time. This means that if a version of a library l_a depends on l_b , and the version that comes after l_a (denoted $\text{succ}(l_a)$) depends on $\text{succ}(l_b)$, we will fail to pick up breaking changes which are due to calling methods which have different semantics *due to the transitive upgrade from l_b to $\text{succ}(l_b)$* . We will, however, pick up breaking changes due to calls to any version of l_b , despite it being a transitive dependency.

Insufficient semantic information. There is also the issue of insufficient semantic information being statically present in source code. An example is the access levels of methods in Python, which are mostly implicit – while the runtime does treat underscore-prefixed names specially in some contexts, it is mostly a convention. Whether or not an API is meant to be internal is not always deducible from source code, and so we necessarily over-approximate, employ heuristics, and guess when analyzing such programs.

6 CONCLUSION AND FUTURE WORK

We presented a static analysis for computing diffs between libraries and determining if an upgrade introduces an API incompatibility. It is scalable, supporting real-time diff queries involving arbitrary pairs of library versions. This makes it lightweight enough to be part of a continuous integration/delivery (CI/CD) pipeline, enabling a vulnerability remediation product that is able to automatically upgrade libraries in 10% of cases. We also evaluated adherence to the semantic versioning scheme on Maven Central, PyPI, and RubyGems, finding that 26% of library versions are in violation.

For future work we would like to improve the accuracy of the analysis and lower false positive rates. A dynamic analysis could be combined with our static call graphs to improve their accuracy. Usability and false negative rates would both be improved by considering more language features in our analysis. Another direction is to improve how actionable results are by suggesting or generating test cases which exercise the code paths in an application which

we know to be involved in a breaking upgrade, or by optimizing the selection of upgrade versions to minimize the manual work users must perform.

REFERENCES

- [1] 2014. Why Semantic Versioning Isn't. <https://gist.github.com/jashkenas/cbd2b088e20279ae2c8e>
- [2] Christopher Bogart, Christian Kästner, James Herbsleb, and Ferdian Thung. 2016. How to Break an API: Cost Negotiation and Community Values in Three Software Ecosystems. In *Proceedings of the 2016 24th ACM SIGSOFT International Symposium on Foundations of Software Engineering (FSE 2016)*. ACM, New York, NY, USA, 109–120. <https://doi.org/10.1145/2950290.2950325>
- [3] Marcel Böhme, Abhik Roychoudhury, and Bruno C. d. S. Oliveira. 2013. Regression Testing of Evolving Programs. *Advances in Computers* 89 (2013), 53–88.
- [4] Barthelemy Dagenais and Martin P. Robillard. 2009. SemDiff: Analysis and Recommendation Support for API Evolution. In *Proceedings of the 31st International Conference on Software Engineering (ICSE '09)*. IEEE Computer Society, Washington, DC, USA, 599–602. <https://doi.org/10.1109/ICSE.2009.5070565>
- [5] Jeffrey Dean, David Grove, and Craig Chambers. 1995. Optimization of object-oriented programs using static class hierarchy analysis. In *European Conference on Object-Oriented Programming*. Springer, 77–101.
- [6] Jean-Rémy Fallier, Floréal Morandat, Xavier Blanc, Matias Martinez, and Martin Monperrus. 2014. Fine-grained and accurate source code differencing. In *Proceedings of the 29th ACM/IEEE international conference on Automated software engineering*. ACM, 313–324.
- [7] Darius Foo, Jason Yeo, Ming Yi Ang, and Asankhaya Sharma. 2018. SGL: A domain-specific language for large-scale analysis of open-source code. *IEEE Cybersecurity Development, SecDev* (2018).
- [8] Johannes Henkel and Amer Diwan. 2005. CatchUp!: Capturing and Replaying Refactorings to Support API Evolution. In *Proceedings of the 27th International Conference on Software Engineering (ICSE '05)*. ACM, New York, NY, USA, 274–283. <https://doi.org/10.1145/1062455.1062512>
- [9] Raula Gaikovina Kula, Daniel M German, Ali Ouni, Takashi Ishio, and Katsuro Inoue. [n. d.]. Do developers update their library dependencies? *Empirical Software Engineering* ([n. d.]), 1–34.
- [10] Gianluca Mezzetti, Anders Møller, and Martin Toldam Torp. [n. d.]. Type Regression Testing to Detect Breaking Changes in Node.js Libraries. ([n. d.]).
- [11] Samim Mirhosseini and Chris Parnin. 2017. Can automated pull requests encourage software developers to upgrade out-of-date dependencies?. In *Proceedings of the 32nd IEEE/ACM International Conference on Automated Software Engineering*. IEEE Press, 84–94.
- [12] Eugene W Myers. 1986. An O(ND) difference algorithm and its variations. *Algorithmica* 1, 1–4 (1986), 251–266.
- [13] S. Raemaekers, A. van Deursen, and J. Visser. 2014. Semantic Versioning versus Breaking Changes: A Study of the Maven Repository. In *2014 IEEE 14th International Working Conference on Source Code Analysis and Manipulation*. 215–224. <https://doi.org/10.1109/SCAM.2014.30>
- [14] Mojtaba Shahin, Muhammad Ali Babar, and Liming Zhu. 2017. Continuous integration, delivery and deployment: a systematic review on approaches, tools, challenges and practices. *IEEE Access* 5 (2017), 3909–3943.
- [15] Vijay Sundarespan, Laurie Hendren, Chrislain Razafimahefa, Raja Vallée-Rai, Patrick Lam, Etienne Gagnon, and Charles Godin. 2000. *Practical virtual method call resolution for Java*. Vol. 35. ACM.
- [16] Mohsen Vakilian, Raluca Sauciu, J David Morgenthaler, and Vahab Mirrokni. 2015. Automated decomposition of build targets. In *Software Engineering (ICSE), 2015 IEEE/ACM 37th IEEE International Conference on*, Vol. 1. IEEE, 123–133.
- [17] Zhenchang Xing and Eleni Stroulia. 2005. UMLDiff: An Algorithm for Object-oriented Design Differencing. In *Proceedings of the 20th IEEE/ACM International Conference on Automated Software Engineering (ASE '05)*. ACM, New York, NY, USA, 54–65. <https://doi.org/10.1145/1101908.1101919>
- [18] Z. Xing and E. Stroulia. 2007. API-Evolution Support with Diff-CatchUp. *IEEE Transactions on Software Engineering* 33, 12 (Dec 2007), 818–836. <https://doi.org/10.1109/TSE.2007.70747>
- [19] Jooyong Yi, Dawei Qi, Shin Hwei Tan, and Abhik Roychoudhury. 2013. Expressing and checking intended changes via software change contracts. In *Proceedings of the 2013 International Symposium on Software Testing and Analysis*. ACM, 1–11.
- [20] Yaqin Zhou and Asankhaya Sharma. 2017. Automated identification of security issues from commit messages and bug reports. In *Proceedings of the 2017 11th Joint Meeting on Foundations of Software Engineering*. ACM, 914–919.